Dimensionality Reduction

Simplifying life

Yordan Darakchiev

Technical Trainer

iordan93@gmail.com

Table of Contents

- sli.do: [#dim-red](#)
- Projections and dimensionality reduction
  - Problem definition
  - PCA
  - Kernel PCA
  - LinDA
- Manifold learning
  - Problem definition
  - Isometric mapping
 - t-SNE
- Real-world examples

Dimensionality Reduction

Reducing the number of features via projection

Principal Component Analysis (PCA)

- A technique for feature extraction
  - Transforms all variables so they are linearly independent
    - Note that this is a linear transformation
- The transformation is done to increase the explained variance in the data
- Number of principal components (PC): $p = \min(m, n)$
  - m – number of features, n – number of observations
  - Sorted from highest to lowest explained variance
- Most widely used practical application
  - Load a high-dimensional dataset with m features
  - Perform PCA, remove last k components
    - Result: dataset with $m - k$ features
    - Note that these **are not** $m - k$ of the original columns!

PCA Intuition

- There is a lot of math (and even physics) involved
- Intuitive explanation: casting a shadow
  - Projection into a lower-dimensional space
    - A nice visualization of PCA
- Example: 2D
  - The first principal component aligns with the axis of most variance in the data
  - The second is perpendicular to the first
    - All components are mutually orthogonal
      - This is one way of saying "linearly independent"
  - Result: The red arrows represent a new coordinate system
    - Also: if we drop the second coordinate, we've achieved dimensionality reduction without too much loss of information

Example: PCA in `scikit-learn`

- Reproduce the 2D example
  - Generate a blob
  - Use a shear transformation to make it look like a rotated ellipse
    - You can look up how to do this, for example in Wikipedia
  - Plot the original data, after the transformation
  - **Rescale the data**
    - This is really important in PCA, otherwise columns with large values will dominate in the PCs
  - Fit a PCA and apply the learned transformation to the original data

```python
from sklearn.decomposition import PCA
pca = PCA().fit(data)
transformed_data = pca.transform(data)
```

  - Plot the transformed data
  - Plot the explained variance ratios of the various components

Kernel PCA

- As we saw, PCA only works for linear transformations
  - We can obtain non-linear transformations using the "kernel trick"
    - Exactly the same way as we did in SVMs
- This allows us to separate non-convex and non-linear data
  - Such as nested circles
  - As before, we need to set the `gamma` parameter
    - Width of the kernel (in the case of `rbf`)
    - If we don't know the exact value, we need to perform grid search
- [Docs](#) (note how linear PCA doesn't separate the data)
- Usage

```python
from sklearn.decomposition import KernelPCA
pca = KernelPCA(kernel = "rbf", gamma = 5).fit(data)
transformed_data = pca.transform(data)
```

Linear Discriminant Analysis (LDA)

- LinDA is a supervised method which tries to identify the attributes that account for the most variance between classes
  - Used in classification cases, with known class labels
  - Returns a transformation of the input data, similar to PCA

- Comparison of 2-component PCA and LinDA on the Iris dataset

```python
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
lda = LinearDiscriminantAnalysis(n_components = 2)
transformed_attributes = lda.fit(attributes, labels).transform(attributes)
```

- Linear Discriminant Analysis and Latent Dirichlet Allocation have the same acronyms, and are both used for dimensionality reduction / transformation
  - They work in completely different ways

Manifold Learning

"Flattening out" surfaces in fewer dimensions

Problem Definition

- Manifold = surface / shape (not necessarily a plane)
- Idea
  - The dimensions are only artificially high
  - Even though there are many features, they can be expressed with a few parameters
    - "Low-dimensional manifold in a high-dimensional space"
    - Non-linear methods
- Example: "swiss roll" dataset
  - 2D "curled" shape in 3D
- Used not only for dimensionality reduction, it's a common way to visualize high-dimensional datasets
  - Especially for classification cases

Swiss Roll in 3D

Isometric Mapping

- Seeks a lower-dimensional projection (embedding) which preserves distances between points
- Algorithm overview
  - Find the k nearest neighbors of each point (kNN)
  - Construct a connectivity graph
    - Two points are connected if they're neighbors (from the first step)
  - Compute shortest paths (Dijkstra, Floyd-Warshall)
  - Perform projection on the graph
    - Similar to PCA
- High-level overview: project a manifold using nearest neighbors

```python
from sklearn.manifold import Isomap
isomap = Isomap(n_neighbors = 5, n_components = 2)
transformed_data = isomap.fit_transform(data)
```

t-SNE

- t-distributed Stochastic Neighbor Embedding
  - Isomap tries to "unfold" a continuous structure
    - This is also true for other algorithms (comparison)
  - t-SNE looks for local clusters in the data
    - Useful for revealing clusters and structure
- Usual implementation: Barnes – Hut (2D or 3D only)
  - To preserve high-dimensional structure, we need to initialize it with PCA (instead of randomly)
  - It can be slow
  - Mainly used for **visualization** in image and text processing

t-SNE embedding of the digits (time 16.68s)

```
from sklearn.manifold import TSNE
tsna = TSNA(n_components = 2, init = "pca")
transformed_data = tsna.fit_transform(data)
```

Real-World Examples

Seeing dimensionality reduction in action

Example: Iris dataset

- Perform PCA on the Iris dataset
  - Use the results to perform feature selection
- Compare the results to isometric mapping and 2D t-SNE
  - Visualize the first 2 components in all three cases
- Choose a classifier and score it before / after PCA
  - E.g. random forest
  - In most cases, first components will give a lower score but the algorithm will perform much faster (fewer variables)
  - Don't forget to transform new incoming data!
    - You can use pipelines to do this
- Compare kNN results on the raw dataset, and on the isomap transformation

Example: Eigenfaces

- Dataset: [Labeled Faces in the Wild](#)
  - Included in `scikit-learn`
- [Example](#)
  - Select only people with many images
    - In this case, $\geq$ 70
  - Use pixel values as raw inputs
  - Perform PCA to select the top 150 components (out of ~3000)
  - Use the PCA components to visualize "eigenfaces"
    - I.e. to see each principal component as an image

```
eigenfaces = pca.components_.reshape((n_components, h, w))
# h, w are image height and width
```

  - Train and optimize an SVM
  - Perform classification

Example: Topics in Text

- Dataset: <u>Newsgroups</u>
  - Included in `scikit-learn`
- <u>Example</u> (LatDA and NMF)
  - LatDA – **Lat**ent **D**irichlet **A**llocation
    - Most widely-used algorithm for topic extraction
  - NMF – **N**on-negative **M**atrix **F**actorization
    - Simpler than LatDA, also used for topic extraction
- Preprocessing: counting words
  - "Bag of words" model
    - Matrix of how often each word occurs in each document
- Visualization of topics: top n words for each topic
- <u>LatDA and k-means clustering on the newsgroups</u>

Summary

- Projections and dimensionality reduction
  - Problem definition
  - PCA
  - Kernel PCA
  - LinDA
- Manifold learning
  - Problem definition
  - Isometric mapping
  - t-SNE
- Real-world examples

Questions?